

3-data-cleaning-and-preparation

February 1, 2024

1 Lab 3: Data Cleaning and Preparation

Objectives: - To be more familiar with Pandas libraries - To gain more hands-on experience in data cleaning and preparation

2 [1] More Reviews on Pandas

1.0) Discover * methods to explore and understand your DataFrame

```
[1]: import pandas as pd

df = pd.read_csv('nss15.csv')
```

C:\Users\woosh\AppData\Local\Temp\ipykernel_18212\1974226225.py:1:

DeprecationWarning:

Pyarrow will become a required dependency of pandas in the next major release of pandas (pandas 3.0),

(to allow more performant data types, such as the Arrow string type, and better interoperability with other libraries)

but was not found to be installed on your system.

If this would cause problems for you,

please provide us feedback at <https://github.com/pandas-dev/pandas/issues/54466>

```
import pandas as pd
```

```
[2]: # see the shape of the dataframe
print(df.shape)
```

(334839, 12)

```
[3]: # seeing the summary of the dataframe
print(df.info())
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 334839 entries, 0 to 334838

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----

```

0  caseNumber      334839 non-null  int64
1  treatmentDate  334839 non-null  object
2  statWeight     334839 non-null  float64
3  stratum        334839 non-null  object
4  age            334839 non-null  int64
5  sex            334837 non-null  object
6  race           205014 non-null  object
7  diagnosis      334839 non-null  int64
8  bodyPart       334839 non-null  int64
9  disposition    334839 non-null  int64
10 location       334839 non-null  int64
11 product       334839 non-null  int64
dtypes: float64(1), int64(7), object(4)
memory usage: 30.7+ MB
None

```

```
[4]: # seeing the stats of the column in dataframe
print(df.describe())
```

	caseNumber	statWeight	age	diagnosis \
count	3.348390e+05	334839.000000	334839.000000	334839.000000
mean	1.510271e+08	39.343028	31.385451	60.154591
std	1.720330e+06	34.142933	26.105098	6.170699
min	1.501032e+08	4.965500	0.000000	41.000000
25%	1.504405e+08	15.059100	10.000000	57.000000
50%	1.507358e+08	15.776200	23.000000	59.000000
75%	1.510231e+08	74.881300	51.000000	64.000000
max	1.603418e+08	97.923900	107.000000	74.000000

	bodyPart	disposition	location	product
count	334839.000000	334839.000000	334839.000000	334839.000000
mean	64.374192	1.307930	2.485451	2098.900854
std	24.002331	0.977627	3.217617	1332.222670
min	0.000000	1.000000	0.000000	106.000000
25%	35.000000	1.000000	0.000000	1211.000000
50%	75.000000	1.000000	1.000000	1807.000000
75%	82.000000	1.000000	5.000000	3265.000000
max	94.000000	9.000000	9.000000	5555.000000

```
[5]: # seeing the first 5 rows of the dataframe
print(df.head())
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race \
0	150733174	7/11/2015	15.7762	V	5	Male	NaN
1	150734723	7/6/2015	83.2157	S	36	Male	White
2	150817487	8/2/2015	74.8813	L	20	Female	NaN
3	150717776	6/26/2015	15.7762	V	61	Male	NaN
4	150721694	7/4/2015	74.8813	L	88	Female	Other

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893

```
[6]: # seeing the last 5 rows of the dataframe
print(df.tail())
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	
334837	150823002	8/8/2015	97.9239	M	38	Female	White	
334838	150723074	6/20/2015	49.2646	M	5	Female	White	

	diagnosis	bodyPart	disposition	location	product
334834	59	76	1	1	1864
334835	68	85	1	0	1931
334836	71	79	1	0	3250
334837	59	82	1	1	464
334838	57	34	1	9	3273

```
[7]: # seeing the list of columns in the dataframe
print(df.columns)
```

```
Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'age', 'sex',
      'race', 'diagnosis', 'bodyPart', 'disposition', 'location', 'product'],
      dtype='object')
```

1.2) Selecting variables * select specific columns from the DataFrame to create a new DataFrame with only those columns

```
[8]: df['age']
```

```
[8]: 0      5
      1     36
      2     20
      3     61
      4     88
      ..
      334834     7
      334835     3
      334836    38
      334837    38
      334838     5
```

Name: age, Length: 334839, dtype: int64

```
[9]: df['age'].head()
```

```
[9]: 0      5
      1     36
      2     20
      3     61
      4     88
      Name: age, dtype: int64
```

```
[10]: df[['caseNumber', 'age']]
```

```
[10]:      caseNumber  age
0      150733174     5
1      150734723    36
2      150817487    20
3      150717776    61
4      150721694    88
...          ...  ...
334834  150739278     7
334835  150733393     3
334836  150819286    38
334837  150823002    38
334838  150723074     5
```

[334839 rows x 2 columns]

```
[11]: # select columns based on the data type
      df.select_dtypes(include=['number'])
```

```
[11]:      caseNumber  statWeight  age  diagnosis  bodyPart  disposition \
0      150733174    15.7762     5         57         33           1
1      150734723    83.2157    36         57         34           1
2      150817487    74.8813    20         71         94           1
3      150717776    15.7762    61         71         35           1
4      150721694    74.8813    88         62         75           1
...          ...  ...  ...  ...  ...  ...
334834  150739278    15.0591     7         59         76           1
334835  150733393     5.6748     3         68         85           1
334836  150819286    15.7762    38         71         79           1
334837  150823002    97.9239    38         59         82           1
334838  150723074    49.2646     5         57         34           1

      location  product
0           9     1267
1           1     1439
```

2	0	3274
3	0	611
4	0	1893
...
334834	1	1864
334835	0	1931
334836	0	3250
334837	1	464
334838	9	3273

[334839 rows x 8 columns]

```
[12]: # select row by .loc
df.loc[0]
```

```
[12]: caseNumber      150733174
      treatmentDate   7/11/2015
      statWeight      15.7762
      stratum         V
      age             5
      sex             Male
      race            NaN
      diagnosis       57
      bodyPart        33
      disposition     1
      location        9
      product         1267
      Name: 0, dtype: object
```

```
[13]: # select column by .loc
df.loc[:, 'treatmentDate': 'diagnosis']
```

```
[13]:  treatmentDate  statWeight  stratum  age  sex  race  diagnosis
0      7/11/2015      15.7762      V    5   Male  NaN        57
1      7/6/2015      83.2157      S   36   Male  White       57
2      8/2/2015      74.8813      L   20  Female  NaN       71
3      6/26/2015      15.7762      V   61   Male  NaN       71
4      7/4/2015      74.8813      L   88  Female  Other       62
5      7/2/2015       5.6748      C    1  Female  White       71
6      6/8/2015      15.7762      V   25   Male  Black       51
```

```
[14]: df.loc[df['age']>80, ['treatmentDate', 'age']]
```

```
[14]:  treatmentDate  age
4      7/4/2015   88
8      7/16/2015  98
39     5/3/2015   88
```

46	4/15/2015	91
63	1/12/2015	97
...
334701	4/27/2015	86
334784	7/7/2015	82
334785	7/11/2015	86
334815	10/28/2015	85
334819	1/13/2015	85

[20422 rows x 2 columns]

```
[15]: # select row by .iloc
df.iloc[0:5]
```

```
[15]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	
1	150734723	7/6/2015	83.2157	S	36	Male	White	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893

```
[16]: # select column by .iloc
df.iloc[:, [0,1,2,3,4]]
```

```
[16]:
```

	caseNumber	treatmentDate	statWeight	stratum	age
0	150733174	7/11/2015	15.7762	V	5
1	150734723	7/6/2015	83.2157	S	36
2	150817487	8/2/2015	74.8813	L	20
3	150717776	6/26/2015	15.7762	V	61
4	150721694	7/4/2015	74.8813	L	88
...
334834	150739278	5/31/2015	15.0591	V	7
334835	150733393	7/11/2015	5.6748	C	3
334836	150819286	7/24/2015	15.7762	V	38
334837	150823002	8/8/2015	97.9239	M	38
334838	150723074	6/20/2015	49.2646	M	5

[334839 rows x 5 columns]

1.3) Filtering the data

```
[17]: # filter rows based on the condition
df[df['age'] > 50]
```

```
[17]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	
7	150704114	6/14/2015	83.2157	S	53	Male	White	
8	150736558	7/16/2015	83.2157	S	98	Male	Black	
16	150901411	8/27/2015	83.2157	S	65	Female	White	
...	
334811	150702215	6/27/2015	15.7762	V	51	Female	NaN	
334815	151100368	10/28/2015	83.2157	S	85	Female	NaN	
334819	150528367	1/13/2015	49.2646	M	85	Female	NaN	
334826	150648619	6/17/2015	15.7762	V	52	Female	White	
334829	150633526	4/4/2015	49.2646	M	51	Female	NaN	

	diagnosis	bodyPart	disposition	location	product
3	71	35	1	0	611
4	62	75	1	0	1893
7	57	30	1	0	5040
8	59	76	1	1	1807
16	59	83	1	1	1817
...
334811	53	83	1	1	1426
334815	57	80	4	1	1807
334819	57	79	5	1	676
334826	64	30	1	1	1842
334829	56	92	1	1	1616

[85235 rows x 12 columns]

```
[18]: # filter coloum based on column name
df.filter(like='age')
```

```
[18]:
```

	age
0	5
1	36
2	20
3	61
4	88
...	...
334834	7
334835	3
334836	38
334837	38
334838	5

[334839 rows x 1 columns]

1.4) Sorting * Sort the DataFrame by its index based on column

```
[19]: # sort the dataframe based on column name and ascending order
df.sort_values(by='statWeight', ascending=False)
```

```
[19]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
67072	150533084	5/15/2015	97.9239	M	89	Male	NaN	
313846	150521217	4/18/2015	97.9239	M	36	Female	NaN	
230135	150857760	8/25/2015	97.9239	M	14	Male	White	
141323	151039262	10/11/2015	97.9239	M	39	Female	White	
230141	150662453	6/5/2015	97.9239	M	11	Female	White	
...	
122009	151146792	11/15/2015	4.9655	C	2	Female	White	
211090	151253201	12/15/2015	4.9655	C	2	Male	White	
317625	160106638	12/25/2015	4.9655	C	1	Male	White	
33679	151256307	12/20/2015	4.9655	C	9	Female	Black	
229596	160148171	12/4/2015	4.9655	C	16	Female	Other	

	diagnosis	bodyPart	disposition	location	product
67072	53	83	1	0	1842
313846	64	79	1	0	5040
230135	64	82	1	8	1807
141323	71	35	1	1	1615
230141	59	88	1	4	3297
...
122009	59	76	1	0	1893
211090	60	88	1	1	661
317625	55	32	1	1	679
33679	57	83	1	0	1842
229596	55	35	1	0	1267

[334839 rows x 12 columns]

```
[20]: # sort the index of the dataframe
df.sort_index()
```

```
[20]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	
1	150734723	7/6/2015	83.2157	S	36	Male	White	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	
...	
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	

334836	150819286	7/24/2015	15.7762	V	38	Male	NaN
334837	150823002	8/8/2015	97.9239	M	38	Female	White
334838	150723074	6/20/2015	49.2646	M	5	Female	White

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893
...
334834	59	76	1	1	1864
334835	68	85	1	0	1931
334836	71	79	1	0	3250
334837	59	82	1	1	464
334838	57	34	1	9	3273

[334839 rows x 12 columns]

1.5) Add/Remove - This section shows how to manipulate the DataFrame's structure

```
[21]: # Dropping the column
df.drop(columns=['disposition'])
```

```
[21]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	
1	150734723	7/6/2015	83.2157	S	36	Male	White	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	
...	
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	
334837	150823002	8/8/2015	97.9239	M	38	Female	White	
334838	150723074	6/20/2015	49.2646	M	5	Female	White	

	diagnosis	bodyPart	location	product
0	57	33	9	1267
1	57	34	1	1439
2	71	94	0	3274
3	71	35	0	611
4	62	75	0	1893
...
334834	59	76	1	1864
334835	68	85	0	1931
334836	71	79	0	3250

334837	59	82	1	464
334838	57	34	9	3273

[334839 rows x 11 columns]

```
[22]: # Adding column and create into a new column
df.assign(new_column=df['diagnosis'] + df['bodyPart'])
```

```
[22]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	
1	150734723	7/6/2015	83.2157	S	36	Male	White	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	
...	
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	
334837	150823002	8/8/2015	97.9239	M	38	Female	White	
334838	150723074	6/20/2015	49.2646	M	5	Female	White	

	diagnosis	bodyPart	disposition	location	product	new_column
0	57	33	1	9	1267	90
1	57	34	1	1	1439	91
2	71	94	1	0	3274	165
3	71	35	1	0	611	106
4	62	75	1	0	1893	137
...
334834	59	76	1	1	1864	135
334835	68	85	1	0	1931	153
334836	71	79	1	0	3250	150
334837	59	82	1	1	464	141
334838	57	34	1	9	3273	91

[334839 rows x 13 columns]

```
[23]: # Removing the column and assigning it to a new variable
# df.pop('age')
```

1.6) Clean missing - to remove rows with missing values or replace missing values with a specified value

```
[24]: # replaceing the missing values with a specified value
df.fillna(value=0)
```

```
[24]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	0	

1	150734723	7/6/2015	83.2157	S	36	Male	White
2	150817487	8/2/2015	74.8813	L	20	Female	0
3	150717776	6/26/2015	15.7762	V	61	Male	0
4	150721694	7/4/2015	74.8813	L	88	Female	Other
...
334834	150739278	5/31/2015	15.0591	V	7	Male	0
334835	150733393	7/11/2015	5.6748	C	3	Female	Black
334836	150819286	7/24/2015	15.7762	V	38	Male	0
334837	150823002	8/8/2015	97.9239	M	38	Female	White
334838	150723074	6/20/2015	49.2646	M	5	Female	White

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893
...
334834	59	76	1	1	1864
334835	68	85	1	0	1931
334836	71	79	1	0	3250
334837	59	82	1	1	464
334838	57	34	1	9	3273

[334839 rows x 12 columns]

```
[25]: # Remove the rows with missing values
df.dropna()
```

```
[25]: caseNumber treatmentDate statWeight stratum age sex race \
1 150734723 7/6/2015 83.2157 S 36 Male White
4 150721694 7/4/2015 74.8813 L 88 Female Other
5 150721815 7/2/2015 5.6748 C 1 Female White
6 150713483 6/8/2015 15.7762 V 25 Male Black
7 150704114 6/14/2015 83.2157 S 53 Male White
...
334830 150628863 6/8/2015 15.7762 V 30 Female White
334831 150607637 5/22/2015 5.6748 C 1 Female Black
334835 150733393 7/11/2015 5.6748 C 3 Female Black
334837 150823002 8/8/2015 97.9239 M 38 Female White
334838 150723074 6/20/2015 49.2646 M 5 Female White

diagnosis bodyPart disposition location product
1 57 34 1 1 1439
4 62 75 1 0 1893
5 71 76 1 1 1715
6 51 33 4 9 1138
```

7	57	30	1	0	5040
...
334830	64	79	1	1	1522
334831	59	94	1	0	1616
334835	68	85	1	0	1931
334837	59	82	1	1	464
334838	57	34	1	9	3273

[205014 rows x 12 columns]

3 [2] Pandas Practice

Now that the knowledge about Pandas is still fresh, let's practice!

2.1) **[Question]** Use pandas to generate a *series* of 20 consecutive numbers, starting from 120.

```
[26]: # write your code here
pd.Series(data = range(120,140), dtype="int")
```

```
[26]: 0    120
      1    121
      2    122
      3    123
      4    124
      5    125
      6    126
      7    127
      8    128
      9    129
     10    130
     11    131
     12    132
     13    133
     14    134
     15    135
     16    136
     17    137
     18    138
     19    139
dtype: int32
```

2.2) **[Question]** Use pandas to generate a *series* of 20 even numbers, starting from 120.

```
[27]: # write your code here
pd.Series(data = range(120,160,2), dtype="int")
```

```
[27]: 0      120
      1      122
      2      124
      3      126
      4      128
      5      130
      6      132
      7      134
      8      136
      9      138
     10      140
     11      142
     12      144
     13      146
     14      148
     15      150
     16      152
     17      154
     18      156
     19      158
dtype: int32
```

2.3) **[Question]** Use pandas to generate a *series* of 50 numbers in the Fibonacci sequence.

(Hint: The Fibonacci sequence is the series of numbers where each number is the sum of the two preceding numbers. For example, 0, 1, 1, 2, 3, 5, ...)

```
[28]: def fibo(n):
      a, b = 0, 1
      for i in range(0, n):
          a, b = b, a + b
      return a
```

```
[29]: # write your code here
      pd.Series(data = [fibo(n) for n in range(50)])
```

```
[29]: 0      0
      1      1
      2      1
      3      2
      4      3
      5      5
      6      8
      7     13
      8     21
      9     34
     10     55
     11     89
```

```
12          144
13          233
14          377
15          610
16          987
17         1597
18         2584
19         4181
20         6765
21        10946
22        17711
23        28657
24        46368
25        75025
26       121393
27       196418
28       317811
29       514229
30       832040
31      1346269
32      2178309
33      3524578
34      5702887
35      9227465
36     14930352
37     24157817
38     39088169
39     63245986
40    102334155
41    165580141
42    267914296
43    433494437
44    701408733
45   1134903170
46   1836311903
47   2971215073
48   4807526976
49   7778742049
dtype: int64
```

2.4) [Question] Use pandas to generate a *series* of 20 random numbers.

```
[32]: # write your code here
import numpy as np
pd.Series(data = np.random.randn(20,))
```

```
[32]: 0      1.152249
      1     -0.591620
      2     -1.791902
      3      0.155772
      4     -0.614648
      5     -0.742754
      6      0.654407
      7     -0.110480
      8     -1.785995
      9     -0.180858
     10      0.442804
     11      0.146879
     12     -1.320452
     13      0.563261
     14      0.862737
     15      0.933455
     16      0.767730
     17      1.461301
     18      0.214919
     19      0.383272
      dtype: float64
```

2.5) **[Question]** Use pandas to generate a *series* of 20 random numbers, indexed in alphabetical order.

```
[33]: # write your code here
      pd.Series(data=np.random.randn(20,), index=[chr(i) for i in range(65,85)])
```

```
[33]: A     -1.281986
      B     -0.578905
      C      1.022104
      D     -1.551667
      E     -0.370524
      F      1.916036
      G      0.543641
      H     -0.304513
      I      0.789417
      J      0.498850
      K     -1.075449
      L      0.554306
      M      0.377492
      N     -2.329619
      O      0.166598
      P      2.139807
      Q     -2.496479
      R     -0.528367
      S      1.819391
      T      0.628939
```

dtype: float64

Next, we're going to use a dataframe which has already been created earlier at the beginning of this notebook. Let's view the first 5 rows (by default).

```
[34]: # df = pd.read_csv('nss15.csv') # uncomment this line if the dataframe has been
      ↪ deleted.
      df.head()
```

```
[34]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	
1	150734723	7/6/2015	83.2157	S	36	Male	White	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893

2.6) [Question] Display the first 12 rows

```
[35]: # write your code here
      df.head(12)
```

```
[35]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	
1	150734723	7/6/2015	83.2157	S	36	Male	White	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	
5	150721815	7/2/2015	5.6748	C	1	Female	White	
6	150713483	6/8/2015	15.7762	V	25	Male	Black	
7	150704114	6/14/2015	83.2157	S	53	Male	White	
8	150736558	7/16/2015	83.2157	S	98	Male	Black	
9	150734928	7/13/2015	74.8813	L	48	Female	Black	
10	150734952	7/4/2015	15.7762	V	20	Male	Black	
11	150821622	7/20/2015	83.2157	S	20	Female	White	

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611

4	62	75	1	0	1893
5	71	76	1	1	1715
6	51	33	4	9	1138
7	57	30	1	0	5040
8	59	76	1	1	1807
9	53	79	1	5	4057
10	59	82	1	1	1894
11	57	36	1	9	1267

2.7) [Question] Display *the last 7 rows*

```
[36]: # write your code here
df.tail(7)
```

```
[36]:      caseNumber treatmentDate  statWeight stratum  age  sex  race \
334832  150747209      7/24/2015    83.2157      S   14  Female  NaN
334833  150747217      7/24/2015    83.2157      S    2   Male  NaN
334834  150739278      5/31/2015   15.0591      V    7   Male  NaN
334835  150733393      7/11/2015    5.6748      C    3  Female  Black
334836  150819286      7/24/2015   15.7762      V   38   Male  NaN
334837  150823002      8/8/2015   97.9239      M   38  Female  White
334838  150723074      6/20/2015   49.2646      M    5  Female  White

      diagnosis  bodyPart  disposition  location  product
334832         62        75           1         5     1807
334833         62        75           1         1     1301
334834         59        76           1         1     1864
334835         68        85           1         0     1931
334836         71        79           1         0     3250
334837         59        82           1         1       464
334838         57        34           1         9     3273
```

2.8) [Question] Display the last 5 rows (by default).

```
[37]: # write your code here
df.tail()
```

```
[37]:      caseNumber treatmentDate  statWeight stratum  age  sex  race \
334834  150739278      5/31/2015   15.0591      V    7   Male  NaN
334835  150733393      7/11/2015    5.6748      C    3  Female  Black
334836  150819286      7/24/2015   15.7762      V   38   Male  NaN
334837  150823002      8/8/2015   97.9239      M   38  Female  White
334838  150723074      6/20/2015   49.2646      M    5  Female  White

      diagnosis  bodyPart  disposition  location  product
334834         59        76           1         1     1864
334835         68        85           1         0     1931
334836         71        79           1         0     3250
```

334837	59	82	1	1	464
334838	57	34	1	9	3273

2.9) **[Question]** Select the column 'statWeight' and display

```
[38]: # write your code here
df['statWeight']
```

```
[38]: 0      15.7762
      1      83.2157
      2      74.8813
      3      15.7762
      4      74.8813
      ...
      334834    15.0591
      334835     5.6748
      334836    15.7762
      334837    97.9239
      334838    49.2646
      Name: statWeight, Length: 334839, dtype: float64
```

2.10) **[Question]** Select the first 20 rows of the column 'statWeight' and display

```
[40]: # write your code here
df['statWeight'].iloc[:20]
df['statWeight'].head(20)
```

```
[40]: 0      15.7762
      1      83.2157
      2      74.8813
      3      15.7762
      4      74.8813
      5       5.6748
      6      15.7762
      7      83.2157
      8      83.2157
      9      74.8813
     10      15.7762
     11      83.2157
     12      15.7762
     13      15.7762
     14      37.6645
     15      83.2157
     16      83.2157
     17       5.6748
     18      15.7762
     19      97.9239
      Name: statWeight, dtype: float64
```

2.11) [Question] Select the last 50 rows of the column 'statWeight' and find/compute the following values: - Minimum - Maximum - Average - Standard Deviation

```
[42]: # write your code here
last50 = df['statWeight'].iloc[-50:]
minn = last50.min()
maxx = last50.max()
mean = last50.mean()
std = last50.std()
last50,minn,maxx,mean,std
```

```
[42]: (334789      5.6748
      334790     83.2157
      334791     74.8813
      334792     74.8813
      334793     97.9239
      334794     15.0591
      334795     15.7762
      334796     74.8813
      334797     15.0591
      334798     49.2646
      334799     15.0591
      334800     15.7762
      334801     49.2646
      334802     74.8813
      334803     74.8813
      334804     74.8813
      334805     15.0591
      334806     97.9239
      334807     15.0591
      334808     15.7762
      334809     15.0591
      334810     97.9239
      334811     15.7762
      334812     85.7374
      334813     97.9239
      334814     85.7374
      334815     83.2157
      334816     15.7762
      334817     15.7762
      334818     97.9239
      334819     49.2646
      334820     15.0591
      334821     15.0591
      334822      5.6748
      334823      5.6748
      334824      5.6748)
```

```

334825    80.8381
334826    15.7762
334827    15.7762
334828    74.8813
334829    49.2646
334830    15.7762
334831     5.6748
334832    83.2157
334833    83.2157
334834    15.0591
334835     5.6748
334836    15.7762
334837    97.9239
334838    49.2646
Name: statWeight, dtype: float64,
5.6748,
97.9239,
45.411078,
34.83805532712222)

```

2.12) [Question] Select the first 25 rows of *two columns* 'statWeight' and 'age', then find/compute the following values for both columns: - Minimum - Maximum - Average - Standard Deviation

```

[44]: # write your code here
last50 = df[['statWeight', 'age']].iloc[-50:]
minn = last50.min()
maxn = last50.max()
mean = last50.mean()
std = last50.std()
print(last50)
minn, maxx, mean, std

```

```

      statWeight  age
334789     5.6748   11
334790    83.2157    9
334791    74.8813   61
334792    74.8813   38
334793    97.9239   13
334794    15.0591    2
334795    15.7762   15
334796    74.8813   46
334797    15.0591   48
334798    49.2646   33
334799    15.0591   31
334800    15.7762   52
334801    49.2646    2
334802    74.8813   29
334803    74.8813   23

```

334804	74.8813	17
334805	15.0591	55
334806	97.9239	18
334807	15.0591	2
334808	15.7762	2
334809	15.0591	4
334810	97.9239	33
334811	15.7762	51
334812	85.7374	3
334813	97.9239	21
334814	85.7374	12
334815	83.2157	85
334816	15.7762	4
334817	15.7762	5
334818	97.9239	5
334819	49.2646	85
334820	15.0591	12
334821	15.0591	27
334822	5.6748	15
334823	5.6748	1
334824	5.6748	1
334825	80.8381	5
334826	15.7762	52
334827	15.7762	8
334828	74.8813	2
334829	49.2646	51
334830	15.7762	30
334831	5.6748	1
334832	83.2157	14
334833	83.2157	2
334834	15.0591	7
334835	5.6748	3
334836	15.7762	38
334837	97.9239	38
334838	49.2646	5

```
[44]: (statWeight    5.6748
      age          1.0000
      dtype: float64,
      97.9239,
      statWeight    45.411078
      age          22.540000
      dtype: float64,
      statWeight    34.838055
      age          22.104769
      dtype: float64)
```

2.13) **[Question]** Select only columns that are of the *type integer*

```
[45]: # write your code here
df.select_dtypes(include=['int'])
```

```
[45]:
```

	caseNumber	age	diagnosis	bodyPart	disposition	location	product
0	150733174	5	57	33	1	9	1267
1	150734723	36	57	34	1	1	1439
2	150817487	20	71	94	1	0	3274
3	150717776	61	71	35	1	0	611
4	150721694	88	62	75	1	0	1893
...
334834	150739278	7	59	76	1	1	1864
334835	150733393	3	68	85	1	0	1931
334836	150819286	38	71	79	1	0	3250
334837	150823002	38	59	82	1	1	464
334838	150723074	5	57	34	1	9	3273

[334839 rows x 7 columns]

2.14) [Question] Select only columns that are of the type *string* or *character*

```
[46]: # write your code here
df.select_dtypes(include=['object'])
```

```
[46]:
```

	treatmentDate	stratum	sex	race
0	7/11/2015	V	Male	NaN
1	7/6/2015	S	Male	White
2	8/2/2015	L	Female	NaN
3	6/26/2015	V	Male	NaN
4	7/4/2015	L	Female	Other
...
334834	5/31/2015	V	Male	NaN
334835	7/11/2015	C	Female	Black
334836	7/24/2015	V	Male	NaN
334837	8/8/2015	M	Female	White
334838	6/20/2015	M	Female	White

[334839 rows x 4 columns]

2.15) [Question] Display only unique values in the column 'race'

```
[47]: # write your code here
df['race'].unique()
```

```
[47]: array([nan, 'White', 'Other', 'Black', 'Asian', 'American Indian'],
      dtype=object)
```

2.16) [Question] Display rows with the following conditions: - Patients are male - The age ranges from 35 to 60 years old - Could be of any race

```
[48]: # write your code here
oldmale = df[(df['sex'] == 'Male') & (df['age'].between(35, 60)) & (~df['race'].
↳isna())]
oldmale
```

```
[48]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
1	150734723	7/6/2015	83.2157	S	36	Male	White	
7	150704114	6/14/2015	83.2157	S	53	Male	White	
32	150908859	8/27/2015	37.6645	L	38	Male	Black	
36	151029422	10/6/2015	97.9239	M	37	Male	White	
44	150407764	3/27/2015	74.8813	L	36	Male	White	
...
334744	151013354	9/30/2015	83.2157	S	49	Male	White	
334751	151042699	10/16/2015	16.5650	V	51	Male	Other	
334769	150648575	6/16/2015	15.7762	V	47	Male	White	
334800	150648581	6/16/2015	15.7762	V	52	Male	White	
334805	150511998	4/20/2015	15.0591	V	55	Male	Black	

	diagnosis	bodyPart	disposition	location	product
1	57	34	1	1	1439
7	57	30	1	0	5040
32	53	36	1	4	5040
36	64	35	1	0	1267
44	62	75	1	0	1842
...
334744	57	82	1	5	1807
334751	59	92	1	0	832
334769	62	75	4	1	1615
334800	64	35	1	1	4074
334805	71	31	6	1	4014

[20683 rows x 12 columns]

2.17) **[Question]** Based on your output in 2.16), select only the columns below to display. - caseNumber - treatmentDate - race - diagnosis - bodyPart - product

```
[49]: # write your code here
oldmale[['caseNumber', 'treatmentDate', 'race', 'diagnosis', 'bodyPart', 'product']]
```

```
[49]:
```

	caseNumber	treatmentDate	race	diagnosis	bodyPart	product
1	150734723	7/6/2015	White	57	34	1439
7	150704114	6/14/2015	White	57	30	5040
32	150908859	8/27/2015	Black	53	36	5040
36	151029422	10/6/2015	White	64	35	1267
44	150407764	3/27/2015	White	62	75	1842
...
334744	151013354	9/30/2015	White	57	82	1807

334751	151042699	10/16/2015	Other	59	92	832
334769	150648575	6/16/2015	White	62	75	1615
334800	150648581	6/16/2015	White	64	35	4074
334805	150511998	4/20/2015	Black	71	31	4014

[20683 rows x 6 columns]

2.18) **[Question]** Let's change the condition a bit. - Patients are female - The age ranges from 5 to 40 years old - Could be of any race

```
[50]: # write your code here
younggirl = df[(df['sex'] == 'Female') & (df['age'].between(5, 40)) &
               (~df['race'].isna())]
younggirl
```

```
[50]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
11	150821622	7/20/2015	83.2157	S	20	Female	White	
13	150666343	6/27/2015	15.7762	V	26	Female	White	
26	151005691	9/29/2015	74.8813	L	27	Female	Black	
43	150413327	3/30/2015	37.6645	L	32	Female	White	
53	150362991	3/11/2015	74.8813	L	13	Female	Black	
...
334817	150639257	6/8/2015	15.7762	V	5	Female	White	
334818	150630000	5/20/2015	97.9239	M	5	Female	Black	
334830	150628863	6/8/2015	15.7762	V	30	Female	White	
334837	150823002	8/8/2015	97.9239	M	38	Female	White	
334838	150723074	6/20/2015	49.2646	M	5	Female	White	

	diagnosis	bodyPart	disposition	location	product
11	57	36	1	9	1267
13	62	75	1	1	1807
26	64	93	1	0	1884
43	71	79	1	0	3216
53	55	30	1	9	1205
...
334817	57	76	2	0	3286
334818	59	94	1	1	1871
334830	64	79	1	1	1522
334837	59	82	1	1	464
334838	57	34	1	9	3273

[45160 rows x 12 columns]

2.19) **[Question]** Likewise, based on your output in 2.18), select only the columns below to display.
- caseNumber - treatmentDate - race - diagnosis - bodyPart - product

```
[51]: # write your code here
younggirl[['caseNumber', 'treatmentDate', 'race', 'diagnosis', 'bodyPart', 'product']]
```



```
[51]:
```

	caseNumber	treatmentDate	race	diagnosis	bodyPart	product
11	150821622	7/20/2015	White	57	36	1267
13	150666343	6/27/2015	White	62	75	1807
26	151005691	9/29/2015	Black	64	93	1884
43	150413327	3/30/2015	White	71	79	3216
53	150362991	3/11/2015	Black	55	30	1205
...
334817	150639257	6/8/2015	White	57	76	3286
334818	150630000	5/20/2015	Black	59	94	1871
334830	150628863	6/8/2015	White	64	79	1522
334837	150823002	8/8/2015	White	59	82	464
334838	150723074	6/20/2015	White	57	34	3273

[45160 rows x 6 columns]

4 [3] Data Cleaning and Preparation

4.0.1 .isnull, .dropna, .fillna

3.1) checking

```
[52]: # isnull checking
df.isnull().sum()
```

```
[52]: caseNumber      0
      treatmentDate  0
      statWeight     0
      stratum        0
      age            0
      sex            2
      race           129825
      diagnosis      0
      bodyPart       0
      disposition    0
      location       0
      product        0
      dtype: int64
```

```
[53]: # percentage of missing values for the race
df.race.isnull().sum()/df.shape[0]*100
```

```
[53]: 38.772365226272925
```

```
[54]: df.shape[0]
```

```
[54]: 334839
```

3.2) Drop column

```
[55]: # remove column by using
df = df.drop(columns=['race'])
```

```
[56]: df.head()
```

```
[56]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	diagnosis	\
0	150733174	7/11/2015	15.7762	V	5	Male	57	
1	150734723	7/6/2015	83.2157	S	36	Male	57	
2	150817487	8/2/2015	74.8813	L	20	Female	71	
3	150717776	6/26/2015	15.7762	V	61	Male	71	
4	150721694	7/4/2015	74.8813	L	88	Female	62	

	bodyPart	disposition	location	product
0	33	1	9	1267
1	34	1	1	1439
2	94	1	0	3274
3	35	1	0	611
4	75	1	0	1893

3.3) Data imputation

```
[57]: # fillna
df['age'] = df['age'].fillna(df['age'].median())
```

3.4) Drop row that have missing value

```
[58]: # remove column by using .dropna()
df = df.dropna()
```

```
[59]: df.isnull().sum()
```

```
[59]: caseNumber      0
      treatmentDate  0
      statWeight     0
      stratum        0
      age            0
      sex            0
      diagnosis      0
      bodyPart       0
      disposition    0
      location       0
      product        0
      dtype: int64
```

4.0.2 Datetime

3.5) Working with the datetime format

```
[60]: df["treatmentDate"] = pd.to_datetime(df["treatmentDate"], format="%m/%d/%Y")
      ↪ #doesn't actually work along the format.
df
```

```
[60]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	diagnosis \
0	150733174	2015-07-11	15.7762	V	5	Male	57
1	150734723	2015-07-06	83.2157	S	36	Male	57
2	150817487	2015-08-02	74.8813	L	20	Female	71
3	150717776	2015-06-26	15.7762	V	61	Male	71
4	150721694	2015-07-04	74.8813	L	88	Female	62
...
334834	150739278	2015-05-31	15.0591	V	7	Male	59
334835	150733393	2015-07-11	5.6748	C	3	Female	68
334836	150819286	2015-07-24	15.7762	V	38	Male	71
334837	150823002	2015-08-08	97.9239	M	38	Female	59
334838	150723074	2015-06-20	49.2646	M	5	Female	57

	bodyPart	disposition	location	product
0	33	1	9	1267
1	34	1	1	1439
2	94	1	0	3274
3	35	1	0	611
4	75	1	0	1893
...
334834	76	1	1	1864
334835	85	1	0	1931
334836	79	1	0	3250
334837	82	1	1	464
334838	34	1	9	3273

[334837 rows x 11 columns]

```
[61]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 334837 entries, 0 to 334838
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   caseNumber      334837 non-null  int64
1   treatmentDate   334837 non-null  datetime64[ns]
2   statWeight      334837 non-null  float64
3   stratum         334837 non-null  object
4   age             334837 non-null  int64
5   sex             334837 non-null  object
6   diagnosis       334837 non-null  int64
7   bodyPart        334837 non-null  int64
```

```

8  disposition      334837 non-null  int64
9  location         334837 non-null  int64
10 product          334837 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(7), object(2)
memory usage: 30.7+ MB

```

```
[62]: df['Year'] = df['treatmentDate'].dt.year
```

```
[63]: df['Month'] = df['treatmentDate'].dt.month
```

```
[64]: df.head()
```

```
[64]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	diagnosis	\
0	150733174	2015-07-11	15.7762	V	5	Male	57	
1	150734723	2015-07-06	83.2157	S	36	Male	57	
2	150817487	2015-08-02	74.8813	L	20	Female	71	
3	150717776	2015-06-26	15.7762	V	61	Male	71	
4	150721694	2015-07-04	74.8813	L	88	Female	62	

	bodyPart	disposition	location	product	Year	Month
0	33	1	9	1267	2015	7
1	34	1	1	1439	2015	7
2	94	1	0	3274	2015	8
3	35	1	0	611	2015	6
4	75	1	0	1893	2015	7

[Question] Can you change the format to DD/MM/YYYY? Show your work.

```
[65]: # write your code here
df['treatmentDate'] = df['treatmentDate'].dt.strftime('%d/%m/%Y')
df
```

```
[65]:
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	diagnosis	\
0	150733174	11/07/2015	15.7762	V	5	Male	57	
1	150734723	06/07/2015	83.2157	S	36	Male	57	
2	150817487	02/08/2015	74.8813	L	20	Female	71	
3	150717776	26/06/2015	15.7762	V	61	Male	71	
4	150721694	04/07/2015	74.8813	L	88	Female	62	
...	
334834	150739278	31/05/2015	15.0591	V	7	Male	59	
334835	150733393	11/07/2015	5.6748	C	3	Female	68	
334836	150819286	24/07/2015	15.7762	V	38	Male	71	
334837	150823002	08/08/2015	97.9239	M	38	Female	59	
334838	150723074	20/06/2015	49.2646	M	5	Female	57	

	bodyPart	disposition	location	product	Year	Month
0	33	1	9	1267	2015	7
1	34	1	1	1439	2015	7

2	94	1	0	3274	2015	8
3	35	1	0	611	2015	6
4	75	1	0	1893	2015	7
...
334834	76	1	1	1864	2015	5
334835	85	1	0	1931	2015	7
334836	79	1	0	3250	2015	7
334837	82	1	1	464	2015	8
334838	34	1	9	3273	2015	6

[334837 rows x 13 columns]

4.0.3 Combine Dataframe by .merge and .concat

3.6 Merge

```
[66]: superstore_order = pd.read_csv('Superstore\superstore_order.csv')
superstore_people = pd.read_csv('Superstore\superstore_people.csv')
superstore_return = pd.read_csv('Superstore\superstore_return.csv')
```

```
<>:1: SyntaxWarning: invalid escape sequence '\s'
<>:2: SyntaxWarning: invalid escape sequence '\s'
<>:3: SyntaxWarning: invalid escape sequence '\s'
<>:1: SyntaxWarning: invalid escape sequence '\s'
<>:2: SyntaxWarning: invalid escape sequence '\s'
<>:3: SyntaxWarning: invalid escape sequence '\s'
C:\Users\woosh\AppData\Local\Temp\ipykernel_18212\1019190268.py:1:
SyntaxWarning: invalid escape sequence '\s'
    superstore_order = pd.read_csv('Superstore\superstore_order.csv')
C:\Users\woosh\AppData\Local\Temp\ipykernel_18212\1019190268.py:2:
SyntaxWarning: invalid escape sequence '\s'
    superstore_people = pd.read_csv('Superstore\superstore_people.csv')
C:\Users\woosh\AppData\Local\Temp\ipykernel_18212\1019190268.py:3:
SyntaxWarning: invalid escape sequence '\s'
    superstore_return = pd.read_csv('Superstore\superstore_return.csv')
```

```
[67]: superstore_order
```

```
[67]:
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	\
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	
4	5	US-2015-108966	11/10/2015	18/10/2015	Standard Class	
...
8875	8876	US-2016-141264	13/08/2016	19/08/2016	Standard Class	
8876	8877	US-2016-141264	13/08/2016	19/08/2016	Standard Class	

8877	8878	CA-2017-126928	17/09/2017	23/09/2017	Standard Class
8878	8879	CA-2017-126928	17/09/2017	23/09/2017	Standard Class
8879	8880	US-2015-107944	23/03/2015	25/03/2015	First Class

	Customer ID	Customer Name	Segment	Country	City \
0	CG-12520	Claire Gute	Consumer	United States	Henderson
1	CG-12520	Claire Gute	Consumer	United States	Henderson
2	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles
3	SO-20335	Sean ODonnell	Consumer	United States	Fort Lauderdale
4	SO-20335	Sean ODonnell	Consumer	United States	Fort Lauderdale
...
8875	CT-11995	Carol Triggs	Consumer	United States	Irving
8876	CT-11995	Carol Triggs	Consumer	United States	Irving
8877	GZ-14470	Gary Zandusky	Consumer	United States	Morristown
8878	GZ-14470	Gary Zandusky	Consumer	United States	Morristown
8879	AM-10360	Alice McCarthy	Corporate	United States	Los Angeles

	...	Postal Code	Region	Product ID	Category	Sub-Category \
0	...	42420	South	FUR-BO-10001798	Furniture	Bookcases
1	...	42420	South	FUR-CH-10000454	Furniture	Chairs
2	...	90036	West	OFF-LA-10000240	Office Supplies	Labels
3	...	33311	South	FUR-TA-10000577	Furniture	Tables
4	...	33311	South	OFF-ST-10000760	Office Supplies	Storage
...
8875	...	75061	Central	OFF-SU-10003505	Office Supplies	Supplies
8876	...	75061	Central	OFF-AP-10002534	Office Supplies	Appliances
8877	...	7960	East	TEC-MA-10004626	Technology	Machines
8878	...	7960	East	OFF-ST-10000615	Office Supplies	Storage
8879	...	90008	West	OFF-PA-10000659	Office Supplies	Paper

	Product Name	Sales	Quantity \
0	Bush Somerset Collection Bookcase	261.9600	2
1	Hon Deluxe Fabric Upholstered Stacking Chairs ...	731.9400	3
2	Self-Adhesive Address Labels for Typewriters b...	14.6200	2
3	Bretford CR4500 Series Slim Rectangular Table	957.5775	5
4	Eldon Fold N Roll Cart System	22.3680	2
...
8875	Premier Electric Letter Opener	185.3760	2
8876	3.6 Cubic Foot Counter Height Office Refrigerator	58.9240	1
8877	Lexmark 20R1285 X6650 Wireless All-in-One Printer	480.0000	4
8878	SimpliFile Personal File Black Granite 15w x 6...	34.0500	3
8879	TOPS Carbonless Receipt Book Four 2-3/4 x 7-1/...	192.7200	11

	Discount	Profit
0	0.00	41.9136
1	0.00	219.5820
2	0.00	6.8714

```

3          0.45 -383.0310
4          0.20   2.5164
...
8875       0.20  -34.7580
8876       0.80 -153.2024
8877       0.00  225.6000
8878       0.00   9.5340
8879       0.00  92.5056

```

[8880 rows x 21 columns]

```
[68]: superstore_order.merge(superstore_return[superstore_return["Returned"]=="Yes"],
    on="Order ID" ,
    how="inner")\
    [["Customer ID", "Returned"]]\
    .drop_duplicates()
```

```
[68]:      Customer ID Returned
0      ZD-21925      Yes
3      TB-21055      Yes
10     JS-15685      Yes
13     LC-16885      Yes
20     BS-11755      Yes
..
688    ED-13885      Yes
689    TS-21205      Yes
696    MF-17665      Yes
702    SH-19975      Yes
705    RB-19435      Yes
```

[222 rows x 2 columns]

[Question] In your opinion, what information that the result above conveys? Ans: Superstore is checking for honest customer that bought their products and have returned the exact products. hahahahahah. The table show customers that returned their product.

More merging...

```
[69]: superstore_order.merge(superstore_return,
    on="Order ID" ,
    how="inner")
```

```
[69]:      Row ID      Order ID  Order Date  Ship Date      Ship Mode \
0      19  CA-2014-143336  27/08/2014  01/09/2014  Second Class
1      20  CA-2014-143336  27/08/2014  01/09/2014  Second Class
2      21  CA-2014-143336  27/08/2014  01/09/2014  Second Class
3      56  CA-2016-111682  17/06/2016  18/06/2016  First Class
4      57  CA-2016-111682  17/06/2016  18/06/2016  First Class
```

..
702	8870	CA-2017-101805	01/12/2017	06/12/2017	Standard Class
703	8871	CA-2017-101805	01/12/2017	06/12/2017	Standard Class
704	8872	CA-2017-101805	01/12/2017	06/12/2017	Standard Class
705	8873	US-2014-105137	10/10/2014	10/10/2014	Same Day
706	8874	US-2014-105137	10/10/2014	10/10/2014	Same Day

	Customer ID	Customer Name	Segment	Country	City \
0	ZD-21925	Zuschuss Donatelli	Consumer	United States	San Francisco
1	ZD-21925	Zuschuss Donatelli	Consumer	United States	San Francisco
2	ZD-21925	Zuschuss Donatelli	Consumer	United States	San Francisco
3	TB-21055	Ted Butterfield	Consumer	United States	Troy
4	TB-21055	Ted Butterfield	Consumer	United States	Troy
..
702	SH-19975	Sally Hughsby	Corporate	United States	Seattle
703	SH-19975	Sally Hughsby	Corporate	United States	Seattle
704	SH-19975	Sally Hughsby	Corporate	United States	Seattle
705	RB-19435	Richard Bierner	Consumer	United States	Columbus
706	RB-19435	Richard Bierner	Consumer	United States	Columbus

	Region	Product ID	Category	Sub-Category \
0	West	OFF-AR-10003056	Office Supplies	Art
1	West	TEC-PH-10001949	Technology	Phones
2	West	OFF-BI-10002215	Office Supplies	Binders
3	East	OFF-ST-10000604	Office Supplies	Storage
4	East	OFF-PA-10001569	Office Supplies	Paper
..
702	West	OFF-BI-10002003	Office Supplies	Binders
703	West	FUR-FU-10000023	Furniture	Furnishings
704	West	OFF-ST-10002756	Office Supplies	Storage
705	East	TEC-MA-10002694	Technology	Machines
706	East	OFF-BI-10002429	Office Supplies	Binders

	Product Name	Sales	Quantity \
0	Newell 341	8.560	2
1	Cisco SPA 501G IP Phone	213.480	3
2	Wilson Jones Hanging View Binder White 1	22.720	4
3	Home/Office Personal File Carts	208.560	6
4	Xerox 232	32.400	5
..
702	Ibico Presentation Index for Binding Systems	15.920	5
703	Eldon Wave Desk Accessories	70.680	12
704	Tennsco Stur-D-Stor Boltless Shelving 5 Shelve...	541.240	4
705	Hewlett-Packard Deskjet F4180 All-in-One Color...	101.994	2
706	Premier Elliptical Ring Binder Black	18.264	2

Discount Profit Returned

0	0.0	2.4824	Yes
1	0.2	16.0110	Yes
2	0.2	7.3840	Yes
3	0.0	52.1400	Yes
4	0.0	15.5520	Yes
..
702	0.2	5.3730	Yes
703	0.0	31.0992	Yes
704	0.0	5.4124	Yes
705	0.7	-71.3958	Yes
706	0.7	-13.3936	Yes

[707 rows x 22 columns]

3.7) Concatenate

```
[70]: pd.concat([superstore_order, superstore_people], axis=1, join='inner')
```

```
[70]:
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID \
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	DV-13045
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335

	Customer Name	Segment	Country	City	...	\
0	Claire Gute	Consumer	United States	Henderson	...	
1	Claire Gute	Consumer	United States	Henderson	...	
2	Darrin Van Huff	Corporate	United States	Los Angeles	...	
3	Sean ODonnell	Consumer	United States	Fort Lauderdale	...	

	Product ID	Category	Sub-Category \
0	FUR-BO-10001798	Furniture	Bookcases
1	FUR-CH-10000454	Furniture	Chairs
2	OFF-LA-10000240	Office Supplies	Labels
3	FUR-TA-10000577	Furniture	Tables

	Product Name	Sales	Quantity \
0	Bush Somerset Collection Bookcase	261.9600	2
1	Hon Deluxe Fabric Upholstered Stacking Chairs ...	731.9400	3
2	Self-Adhesive Address Labels for Typewriters b...	14.6200	2
3	Bretford CR4500 Series Slim Rectangular Table	957.5775	5

	Discount	Profit	Person	Region
0	0.00	41.9136	Anna Andreadi	West
1	0.00	219.5820	Chuck Magee	East
2	0.00	6.8714	Kelly Williams	Central
3	0.45	-383.0310	Cassandra Brandow	South

[4 rows x 23 columns]

4.0.4 Groupby

```
[71]: superstore_order.groupby(['Segment', 'Ship_
    ↪Mode'])[['Sales', 'Quantity', 'Discount', 'Profit']].sum()
```

```
[71]:
```

		Sales	Quantity	Discount	Profit
Segment	Ship Mode				
Consumer	First Class	138594.9328	2455	110.29	18953.7264
	Same Day	53660.6340	1001	43.85	8555.7193
	Second Class	203605.6822	3489	127.29	24701.9148
	Standard Class	627061.3262	10430	443.05	68864.9892
Corporate	First Class	97720.1209	1670	73.07	12660.2526
	Same Day	41716.5550	366	14.50	1120.9222
	Second Class	130759.9288	2027	71.47	15582.1762
	Standard Class	359359.2109	6203	262.82	49832.6780
Home Office	First Class	76743.8674	924	39.82	11829.8821
	Same Day	20968.5170	343	12.50	3909.3442
	Second Class	77175.1080	1148	37.80	12785.8953
	Standard Class	218325.9795	3595	142.14	27298.5786

[Question] Briefly describe an information that the result above conveys?

Ans: This table is a summary of the sales in all segments. In each segment is divided into difference ship mode. As can be seen, the majority sales quantity is from consumer segment. The profit significantly come from standard class ship mode from every segment. Same day ship mode is not popular like the others.

```
[72]: superstore_order["Profit Ratio"] = superstore_order["Profit"]/
    ↪superstore_order["Sales"]
```

```
[73]: superstore_order.groupby(["Category", "Sub-Category"]).agg(mean_profit_ratio =_
    ↪("Profit Ratio", "mean"))
```

```
[73]:
```

		mean_profit_ratio
Category	Sub-Category	
Furniture	Bookcases	-0.127756
	Chairs	0.045028
	Furnishings	0.140782
	Tables	-0.147916
Office Supplies	Appliances	-0.145513
	Art	0.251678
	Binders	-0.191641
	Envelopes	0.421913
	Fasteners	0.301157

	Labels	0.429984
	Paper	0.425586
	Storage	0.092382
	Supplies	0.104970
Technology	Accessories	0.219012
	Copiers	0.317826
	Machines	-0.059535
	Phones	0.118926

[Question] Briefly describe an information that the result above conveys? Ans: The table show the mean of profit ratio to see which product can make a profit. As can be seen roughly, office supplies is the most profitable category. Labels is the best one that can make a profit comparing to it's sales.

4.0.5 Pivot and Melt

Pivot

```
[74]: superstore_order.pivot_table(index="State", columns="Ship Mode", values="Order_ID",
    ↪aggfunc="count").fillna(0).head(10)
```

```
[74]: Ship Mode      First Class  Same Day  Second Class  Standard Class
State
Alabama           9.0         1.0         18.0         30.0
Arizona          42.0        15.0         22.0        123.0
Arkansas          10.0         2.0         8.0         35.0
California       302.0       106.0       346.0       1000.0
Colorado          43.0         5.0        32.0         95.0
Connecticut       19.0         8.0        11.0         39.0
Delaware          16.0         2.0        13.0         55.0
District of Columbia  0.0         0.0         3.0          7.0
Florida           47.0        25.0        57.0        210.0
Georgia           19.0        15.0        31.0        108.0
```

```
[75]: pivot_table_result = superstore_order.pivot_table(index="State", columns="Ship_
    ↪Mode", values="Order ID", aggfunc="count").fillna(0)
print(pivot_table_result)
```

```
Ship Mode      First Class  Same Day  Second Class  Standard Class
State
Alabama           9.0         1.0         18.0         30.0
Arizona          42.0        15.0         22.0        123.0
Arkansas          10.0         2.0         8.0         35.0
California       302.0       106.0       346.0       1000.0
Colorado          43.0         5.0        32.0         95.0
Connecticut       19.0         8.0        11.0         39.0
Delaware          16.0         2.0        13.0         55.0
District of Columbia  0.0         0.0         3.0          7.0
```

Florida	47.0	25.0	57.0	210.0
Georgia	19.0	15.0	31.0	108.0
Idaho	3.0	0.0	2.0	13.0
Illinois	58.0	24.0	96.0	249.0
Indiana	13.0	3.0	30.0	79.0
Iowa	1.0	1.0	4.0	17.0
Kansas	6.0	1.0	2.0	15.0
Kentucky	12.0	5.0	49.0	62.0
Louisiana	7.0	2.0	14.0	15.0
Maine	0.0	0.0	0.0	5.0
Maryland	18.0	7.0	12.0	63.0
Massachusetts	14.0	4.0	35.0	71.0
Michigan	20.0	16.0	43.0	151.0
Minnesota	9.0	4.0	13.0	59.0
Mississippi	3.0	4.0	7.0	36.0
Missouri	7.0	2.0	20.0	24.0
Montana	1.0	1.0	0.0	13.0
Nebraska	6.0	3.0	6.0	20.0
Nevada	4.0	1.0	12.0	17.0
New Hampshire	2.0	0.0	10.0	13.0
New Jersey	5.0	1.0	20.0	87.0
New Mexico	1.0	0.0	9.0	22.0
New York	155.0	57.0	183.0	606.0
North Carolina	36.0	14.0	40.0	139.0
North Dakota	0.0	0.0	5.0	2.0
Ohio	66.0	47.0	84.0	199.0
Oklahoma	5.0	6.0	7.0	44.0
Oregon	20.0	0.0	15.0	81.0
Pennsylvania	103.0	9.0	78.0	341.0
Rhode Island	16.0	0.0	21.0	16.0
South Carolina	3.0	5.0	18.0	16.0
South Dakota	2.0	0.0	0.0	9.0
Tennessee	21.0	2.0	24.0	118.0
Texas	125.0	37.0	161.0	537.0
Utah	4.0	2.0	19.0	28.0
Vermont	0.0	0.0	1.0	2.0
Virginia	39.0	4.0	33.0	115.0
Washington	56.0	34.0	97.0	265.0
West Virginia	0.0	0.0	0.0	3.0
Wisconsin	12.0	3.0	10.0	66.0
Wyoming	0.0	0.0	0.0	1.0

Melt

```
[76]: melted_result = pd.melt(pivot_table_result.reset_index(), id_vars=["State"],
    ↪var_name="Ship Mode", value_name="Order Count")
print(melted_result)
```

	State	Ship Mode	Order Count
0	Alabama	First Class	9.0
1	Arizona	First Class	42.0
2	Arkansas	First Class	10.0
3	California	First Class	302.0
4	Colorado	First Class	43.0
..
191	Virginia	Standard Class	115.0
192	Washington	Standard Class	265.0
193	West Virginia	Standard Class	3.0
194	Wisconsin	Standard Class	66.0
195	Wyoming	Standard Class	1.0

[196 rows x 3 columns]

5 [4] Some more questions!

Let's practice more using the superstore dataset :D

4.1) [Question] From the superstore_order, display the ascending order considering the average values of the 'Profit' column to group the 'Category'.

```
[77]: #enter your code here
print(superstore_order.groupby('Category')['Profit'].mean().
      ↪sort_values(ascending=True))
superstore_order.groupby('Category').agg(Average_profit = ("Profit", "mean"))
```

```
Category
Furniture      8.967320
Office Supplies 19.743848
Technology     81.347862
Name: Profit, dtype: float64
```

```
[77]:          Average_profit
Category
Furniture      8.967320
Office Supplies 19.743848
Technology     81.347862
```

4.2) [Question] Create a new column that calculates the total price (sale*quantity) before discount then group by 'product id' and 'category', then show the mean of the total price

```
[78]: #enter your code here
superstore_order["total_price"] =_
      ↪superstore_order["Sales"]*superstore_order["Quantity"]
superstore_order.groupby(['Product ID', 'Category']).agg(meam_total_price =_
      ↪("total_price", "mean"))
```

```

[78]:
      Product ID      Category      meam_total_price
FUR-B0-10000112  Furniture      7426.566000
FUR-B0-10000330  Furniture      1258.192000
FUR-B0-10000362  Furniture      1726.898000
FUR-B0-10000468  Furniture       426.532400
FUR-B0-10000711  Furniture      3194.100000
...
TEC-PH-10004912  Technology      747.320000
TEC-PH-10004922  Technology      673.249500
TEC-PH-10004924  Technology       57.149333
TEC-PH-10004959  Technology      412.009000
TEC-PH-10004977  Technology     2441.475429

[1846 rows x 1 columns]

```