

classification-toy-example-copy

March 9, 2024

1 Lab 3: Introducing Classification

Objectives: - To gain hands-on experience classifying small dataset - To implement concepts related to Decision Tree classifier (i.e. Entropy, Information Gain), along with the Decision Tree algorithm

```
[ ]: import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
# Read the data
df = pd.read_csv('toy_data.csv')
df
```

```
[ ]: print(df.info())
```

```
[ ]: df_num = df.copy()
d = {'<=30':0, '31-40':1, '>40':2}
df_num['age'] = df_num['age'].map(d)
d = {'high':2, 'medium':1, 'low':0}
df_num['income'] = df_num['income'].map(d)
d = {'yes':1, 'no':0}
df_num['student'] = df_num['student'].map(d)
df_num['buys computer'] = df_num['buys computer'].map(d)
d = {'excellent':1, 'fair':0}
df_num['credit rating'] = df_num['credit rating'].map(d)
df_num
```

2 Calculate Gain

2.0.1 entropy(t)

```
[ ]: entropy_t = 0
n=df_num.shape[0]
countt = df_num['buys computer'].value_counts()
entropy_t = -countt[1]*math.log2(countt[1]/n)/n - countt[0]*math.log2(countt[0]/
↪n)/n
print(entropy_t)
```

2.0.2 entropy(i)

```
[ ]: def entropy (df, target_att):
    ans = 0
    # print(f"n = {n}")
    for data in df[target_att].unique():
        print("data",data)
        n_i = df[(df[target_att] == data)].shape[0]
        # print("size",n_i)
        p_jyes = df[(df[target_att] == data) & (df['buys computer'] == 1)].
        ↪shape[0]/n_i
        # print(f"p(j=yes/{data})={p_jyes}")
        p_jno = df[(df[target_att] == data) & (df['buys computer'] == 0)].
        ↪shape[0]/n_i
        # print(f"p(j=no/{data})={p_jno}")
        if (p_jyes == 0 or p_jno == 0):
            log_j = 0
        else:
            log_j = -(p_jyes * math.log2(p_jyes) + p_jno * math.log2(p_jno))
        # print(f"log(j/{data})={log_j}")
        ans += (n_i*log_j/n)
        # print("\n")
    return ans
```

2.0.3 Gain fn

```
[ ]: def gain_fn (target_att):
    return entropy_t - entropy(df_num,target_att)
```

3 Calculation

```
[ ]: gain_fn("age")
```

```
[ ]: gain_fn("income")
```

```
[ ]: gain_fn("student")
```

```
[ ]: gain_fn("credit rating")
```

4 Decision Tree

make it all numerical

```
[ ]: features = ['age','income','student','credit rating']
X = df_num[features]
y = df_num['buys computer']
```

```
print(X)
print(y)
```

```
[ ]: print(pd.concat([X,y],axis=1))
```

```
[ ]: class Node():
    """
    A class representing a node in a decision tree.
    """

    def __init__(self, feature=None, threshold=None, left=None, right=None,
    ↪gain=None, value=None):
        """
        Initializes a new instance of the Node class.

        Args:
            feature: The feature used for splitting at this node. Defaults to
            ↪None.
            threshold: The threshold used for splitting at this node. Defaults
            ↪to None.
            left: The left child node. Defaults to None.
            right: The right child node. Defaults to None.
            gain: The gain of the split. Defaults to None.
            value: If this node is a leaf node, this attribute represents the
            ↪predicted value
                   for the target variable. Defaults to None.
        """
        self.feature = feature
        self.threshold = threshold
        self.left = left
        self.right = right
        self.gain = gain
        self.value = value
```

```
[70]: class DecisionTree():
    """
    A decision tree classifier for binary classification problems.
    """

    def __init__(self, min_samples=2, max_depth=2):
        """
        Constructor for DecisionTree class.

        Parameters:
            min_samples (int): Minimum number of samples required to split an
            ↪internal node.
```

```

        max_depth (int): Maximum depth of the decision tree.
    """
    self.min_samples = min_samples
    self.max_depth = max_depth

def split_data(self, dataset, feature, threshold):
    """
    Splits the given dataset into two datasets based on the given feature
    and threshold.

    Parameters:
        dataset (ndarray): Input dataset.
        feature (int): Index of the feature to be split on.
        threshold (float): Threshold value to split the feature on.

    Returns:
        left_dataset (ndarray): Subset of the dataset with values less than
        or equal to the threshold.
        right_dataset (ndarray): Subset of the dataset with values greater
        than the threshold.
    """
    # Create empty arrays to store the left and right datasets
    left_dataset = []
    right_dataset = []

    # Loop over each row in the dataset and split based on the given
    feature and threshold
    for row in dataset:
        if row[feature] <= threshold:
            left_dataset.append(row)
        else:
            right_dataset.append(row)

    # Convert the left and right datasets to numpy arrays and return
    left_dataset = np.array(left_dataset)
    right_dataset = np.array(right_dataset)
    return left_dataset, right_dataset

def entropy(self, y):
    """
    Computes the entropy of the given label values.

    Parameters:
        y (ndarray): Input label values.

    Returns:
        entropy (float): Entropy of the given label values.
    """

```

```

"""
entropy = 0

# Find the unique label values in y and loop over each value
# THIS IS FUCKING CLEVERRRR!!
# AND IT'S THE SAME AS PROF SAID
labels = np.unique(y)
for label in labels:
    # Find the examples in y that have the current label
    label_examples = y[y == label]
    # Calculate the ratio of the current label in y
    p1 = len(label_examples) / len(y)
    # Calculate the entropy using the current label and ratio
    entropy += -p1 * np.log2(p1)

# Return the final entropy value
return entropy

def information_gain(self, parent, left, right):
    """
    Computes the information gain from splitting the parent dataset into
    two datasets.

    Parameters:
    parent (ndarray): Input parent dataset.
    left (ndarray): Subset of the parent dataset after split on a
    feature.
    right (ndarray): Subset of the parent dataset after split on a
    feature.

    Returns:
    information_gain (float): Information gain of the split.
    """

    # set initial information gain to 0
    information_gain = 0
    # compute entropy for parent
    parent_entropy = self.entropy(parent)
    # calculate weight for left and right nodes
    weight_left = len(left) / len(parent)
    weight_right = len(right) / len(parent)
    # compute entropy for left and right nodes
    entropy_left, entropy_right = self.entropy(left), self.entropy(right)
    # calculate weighted entropy
    weighted_entropy = weight_left * entropy_left + weight_right *
    entropy_right
    # calculate information gain
    information_gain = parent_entropy - weighted_entropy

```

```

        return information_gain

def best_split(self, dataset, num_samples, num_features):
    """
    Finds the best split for the given dataset.

    Args:
        dataset (ndarray): The dataset to split.
        num_samples (int): The number of samples in the dataset.
        num_features (int): The number of features in the dataset.

    Returns:
        dict: A dictionary with the best split feature index, threshold, gain,
            left and right datasets.
    """
    # dictionary to store the best split values
    best_split = {'gain': -1, 'feature': None, 'threshold': None}
    # loop over all the features
    for feature_index in range(num_features):
        # get the feature at the current feature_index
        feature_values = dataset[:, feature_index]
        # get unique values of that feature
        thresholds = np.unique(feature_values)
        # loop over all values of the feature
        for threshold in thresholds:
            # get left and right datasets
            left_dataset, right_dataset = self.split_data(dataset,
↪feature_index, threshold)
            # check if either datasets is empty
            if len(left_dataset) and len(right_dataset):
                # get y values of the parent and left, right nodes
                y, left_y, right_y = dataset[:, -1], left_dataset[:, -1],
↪right_dataset[:, -1]
                # compute information gain based on the y values
                information_gain = self.information_gain(y, left_y, right_y)
                # update the best split if conditions are met
                if information_gain > best_split["gain"]:
                    best_split["feature"] = feature_index
                    best_split["threshold"] = threshold
                    best_split["left_dataset"] = left_dataset
                    best_split["right_dataset"] = right_dataset
                    best_split["gain"] = information_gain
    return best_split

def calculate_leaf_value(self, y):

```

```

"""
Calculates the most occurring value in the given list of y values.

Args:
    y (list): The list of y values.

Returns:
    The most occurring value in the list.
"""
y = list(y)
#get the highest present class in the array
most_occuring_value = max(y, key=y.count)
return most_occuring_value

def build_tree(self, dataset, current_depth=0):
    """
Recursively builds a decision tree from the given dataset.

Args:
    dataset (ndarray): The dataset to build the tree from.
    current_depth (int): The current depth of the tree.

Returns:
    Node: The root node of the built decision tree.
    """

    # split the dataset into X, y values
    X, y = dataset[:, :-1], dataset[:, -1]
    n_samples, n_features = X.shape
    # keeps splitting until stopping conditions are met
    if n_samples >= self.min_samples and current_depth <= self.max_depth:
        # Get the best split
        best_split = self.best_split(dataset, n_samples, n_features)
        # Check if gain isn't zero
        if best_split["gain"]:
            # continue splitting the left and the right child. Increment
↪ current depth
            left_node = self.build_tree(best_split["left_dataset"], ↪
↪ current_depth + 1)
            right_node = self.build_tree(best_split["right_dataset"], ↪
↪ current_depth + 1)
            # return decision node
            return Node(best_split["feature"], best_split["threshold"],
                        left_node, right_node, best_split["gain"])

    # compute leaf node value
    leaf_value = self.calculate_leaf_value(y)
    # return leaf node value

```

```

    return Node(value=leaf_value)

def fit(self, X, y):
    """
    Builds and fits the decision tree to the given X and y values.

    Args:
    X (ndarray): The feature matrix.
    y (ndarray): The target values.
    """
    # print(X.shape, y.shape)
    # dataset = np.concatenate((X, y), axis=1)
    dataset = pd.concat([X,y],axis=1)
    dataset = dataset.to_numpy()
    self.root = self.build_tree(dataset)

def predict(self, X):
    """
    Predicts the class labels for each instance in the feature matrix X.

    Args:
    X (ndarray): The feature matrix to make predictions for.

    Returns:
    list: A list of predicted class labels.
    """
    # Create an empty list to store the predictions
    predictions = []
    X= X.to_numpy()
    # For each instance in X, make a prediction by traversing the tree
    for x in X:
        prediction = self.make_prediction(x, self.root)
        # Append the prediction to the list of predictions
        predictions.append(prediction)
    # Convert the list to a numpy array and return it
    predictions = np.array(predictions)
    return predictions

def make_prediction(self, x, node):
    """
    Traverses the decision tree to predict the target value for the given
    ↪ feature vector.

    Args:
    x (ndarray): The feature vector to predict the target value for.
    node (Node): The current node being evaluated.

```



```

Returns:
The predicted target value for the given feature vector.
"""

# if the node has value i.e it's a leaf node extract it's value
if node.value != None:
    return node.value
else:
    #if it's node a leaf node we'll get it's feature and traverse
    ↪ through the tree accordingly
    feature = x[node.feature]
    if feature <= node.threshold:
        return self.make_prediction(x, node.left)
    else:
        return self.make_prediction(x, node.right)

```

```

[71]: def train_test_split(X, y, test_size=0.2):
    """
    Splits data into training and testing sets without shuffling.

    Args:
    X (ndarray): The feature matrix.
    y (ndarray): The target values.
    test_size (float, optional): The proportion of data to be used for the
    ↪ test set. Defaults to 0.2.

    Returns:
    tuple: A tuple containing the training and testing data (X_train, X_test,
    ↪ y_train, y_test).
    """

    if not isinstance(test_size, (float, int)):
        raise ValueError("test_size should be a float or an integer")
    elif test_size < 0 or test_size > 1:
        raise ValueError("test_size should be between 0.0 and 1.0")

    n_samples = X.shape[0]
    test_index = int(n_samples * test_size)

    X_train, X_test = X[:test_index], X[test_index:]
    y_train, y_test = y[:test_index], y[test_index:]

    return X_train, X_test, y_train, y_test

```

```

[72]: def accuracy(y_true, y_pred):
    """
    Computes the accuracy of a classification model.

```

Parameters:

y_true (numpy array): A numpy array of true labels for each data point.

y_pred (numpy array): A numpy array of predicted labels for each data_
point.

Returns:

float: The accuracy of the model

"""

y_true = y_true.flatten()

total_samples = len(y_true)

correct_predictions = np.sum(y_true == y_pred)

return (correct_predictions / total_samples)

```
[73]: def balanced_accuracy(y_true, y_pred):  
    """Calculate the balanced accuracy for a multi-class classification problem.
```

Parameters

y_true (numpy array): A numpy array of true labels for each data point.

y_pred (numpy array): A numpy array of predicted labels for each data_
point.

Returns

balanced_acc : The balanced accuracy of the model

"""

y_pred = np.array(y_pred)

y_true = y_true.flatten()

Get the number of classes

n_classes = len(np.unique(y_true))

Initialize an array to store the sensitivity and specificity for each_

class

sen = []

spec = []

Loop over each class

for i in range(n_classes):

Create a mask for the true and predicted values for class i

mask_true = y_true == i

mask_pred = y_pred == i

Calculate the true positive, true negative, false positive, and false_

negative values

TP = np.sum(mask_true & mask_pred)

```

    TN = np.sum((mask_true != True) & (mask_pred != True))
    FP = np.sum((mask_true != True) & mask_pred)
    FN = np.sum(mask_true & (mask_pred != True))

    # Calculate the sensitivity (true positive rate) and specificity (true
    ↪negative rate)
    sensitivity = TP / (TP + FN)
    specificity = TN / (TN + FP)

    # Store the sensitivity and specificity for class i
    sen.append(sensitivity)
    spec.append(specificity)

    # Calculate the balanced accuracy as the average of the sensitivity and
    ↪specificity for each class
    average_sen = np.mean(sen)
    average_spec = np.mean(spec)
    balanced_acc = (average_sen + average_spec) / n_classes

    return balanced_acc

```

```

[74]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.8)
print(X_train)
# print(X_test)
print(y_train)
# print(y_test)

```

	age	income	student	credit	rating
0	0	2	0		0
1	0	2	0		1
2	1	2	0		0
3	2	1	0		0
4	2	0	1		0
5	2	0	1		1
6	1	0	1		1
7	0	1	0		0
8	0	0	1		0
9	2	1	1		0
10	0	1	1		1
0	0				
1	0				
2	1				
3	1				
4	1				
5	0				
6	1				
7	0				
8	1				

```
9      1
10     1
Name: buys computer, dtype: int64
```

```
[78]: #create model instance
model = DecisionTree(2,10)

# Fit the decision tree model to the training data.
model.fit(X_train, y_train)

# Use the trained model to make predictions on the test data.
predictions = model.predict(X_test)

# Calculate evaluating metrics
print(f"Model's Accuracy: {accuracy(y_test.to_numpy(), predictions)}")
print(f"Model's Balanced Accuracy: {balanced_accuracy(y_test.to_numpy(),
↪ predictions)}")
```

```
Model's Accuracy: 1.0
Model's Balanced Accuracy: 1.0
```

```
[ ]:
```